

Due: May 4th (Saturday), 6pm, submit to courseworks.

You may use “standard” arguments without a formal proof—either cite from class or add one sentence showing the main idea.

Problem 1: Independence properties of Tabulation Hashing

Recall the definition of Tabulation hashing: Each w -bit key $x = x_1 \dots x_w$ is partitioned into c blocks of $B = (w/c)$ bits (characters) each $x'_1 \dots x'_c$, and hashed into the XOR of c totally random hash tables : $T(x) := T_1[x'_1] \oplus \dots T_c[x'_c]$, where each T_i is a random table $\in_R [m]^{2^B}$.

1. Prove that this hash function is 3-wise independent, that is: for any three keys x, y, z , the tuple $(T(x), T(y), T(z))$ is equally likely to get mapped to any value in $[m]^3$.
2. Prove that Tabulation hashing is not 4-wise independent: show that there are key values x, y, w, z such that their hashes $T(x), T(y), T(w), T(z)$ are not independent.

Problem 2: Cuckoo Hashing

Prove that $\Pr[\text{Insert}(x) \text{ to Cuckoo hash traverses a } k\text{-length simple path}] \geq 2^{-O(k)}$, (when we hash n keys to 2 completely random tables g, h each of size $m = (1 + \epsilon)n$). This is tight by the argument proved in class (we showed the probability of this event is $\leq 2^{-\Omega(k)}$).

Also show that the probability that cuckoo hashing fails (i.e. it makes $\Omega(\log n)$ “hops” when inserting some key) is $\Omega(1/n)$ (hint: what is the simplest case in which such event happens? Obviously we have excluded graphs with long chains, what other cases are there?).

Problem 3: Longest Non-Overlapping Substring

Devise a linear-time algorithm that for a given string S finds the longest substring x such that there are two positions i, j : $S_i S_{i+1} \dots S_{i+|x|-1} = S_j S_{j+1} \dots S_{j+|x|-1} = x$ and $i + |x| - 1 < j$, i.e. substring x occurs at least twice in S and these two occurrences don't overlap.

Problem 4: Longest Common Substring in linear time

The LCS problem is, given 2 strings a, b , to find the longest (contiguous) common substring of T_1, T_2 . (e.g., if $T_1 = \text{“alibabalc”}$ and $T_2 = \text{“abac”}$, then $LCS(T_1, T_2) = \text{“aba”}$). Use suffix trees to solve this algorithmic problem in $O(|T_1| + |T_2|)$ time.

Problem 5: Least Common Ancestor

Consider the (static) LCA problem: Given a binary tree T on n nodes, preprocess it so that for every pair of vertices $u, v \in T$, we can quickly compute $LCA(u, v)$, that is, the node v of minimum depth in T , which is an ancestor of both u, v .

1. Consider the following “Range Minimum” problem RMQ_n : Preprocess an array of length n so that for any 2 entries $i, j \in [n]$, the DS must output the minimum element $x \in A[i] \dots A[j]$. Show that the LCA problem reduces to RMQ_n , i.e., that any (s, t) -DS for RMQ induces a similar DS for LCA.
(Hint: *Euler Tour* of a graph may help here).
2. Show how to solve RMQ_n with $O(n \lg n)$ words of space and constant $t = O(1)$ query time.
(Hint: Consider storing, for every entry in the array, the minimum element at distance 2^i from that entry, for every i . How much space does this take? And why does it suffice for RMQ queries?).
3. For extra credit, show how to reduce the space to linear $O(n)$. To do this, use Indirection + the important fact that the the resulting RMQ problem has all consecutive entries of the array $A[i]$ differ from the next/previous entry $A[i + 1]$ by at most ± 1 .
4. As an application, show that the LCP problem discussed in class (Given a text T and 2 indexes i, j , determine the longest common prefix $LCP(T[i:], T[j:])$) can be solved in $O(n)$ space and $O(1)$ time.

Problem 6: Finding Blobs

Consider the following dynamic problem: given a changing forest T with n vertices, each either black or white, you are asked to compute sizes of the “blobs” around some vertices under edge addition/removal and color changes. A “blob” around vertex v is the largest connected component containing v that consists of vertices of the same color.

More specifically, your algorithm should support the following types of queries:

1. $\text{Change}(v)$ — change color of vertex v ,
2. $\text{BlobSize}(v)$ — return the size of the “blob” around vertex v ,
3. $\text{Add}(u, v)$ — add an edge between u and v ,
4. $\text{Remove}(u, v)$ — remove an edge between u and v .

It is promised that at every point of time T is a forest. Design an algorithm for that problem with linear space, give bounds on its running time. You will get partial credit if your algorithm supports only Change and BlobSize queries.