

Lecture #9: Cell Probe LBs for Dynamic Range Counting

Instructor: *Omri Weinstein*Scribes: *Weston Jackson*

1 Last Time

Orthogonal Range Counting (ORC): return all points (or sum of weights of all points) in a rectangle R :


$$\sum_{(x,y) \in R} w_{xy}$$

Previously, we used layered range trees to solve in time:

- Static: $t_u = t_q = O(\log^{d-1} n)$
- Dynamic: $t_u = t_q = O(\log^d n)$

For the static case, we used fractional cascading in the last layer to save one log factor. We pay $\log n$ overhead for dynamization, which adds this factor back. However, it is possible to do better $t_u = O(\log^d n), t_q = O(\log^{d-1} n)$ with weight-balanced trees¹.

1D-ORC: Updates are insertions into the number line. Queries ask us to report all points in some interval $[i, j]$:



2 Partial Sums (PS_n) Problem

2.1 Static Case

The PS_n problem is equivalent. Our updates just set the index $A[i] \leftarrow \{0, 1\}$. For queries, we define $PREFIX()$ function as the following:

$$PREFIX(i) = \sum_{j \leq i} A[j]$$

For $QUERY(i, j)$, we make two $PREFIX()$ calls, and just return:

$$PREFIX(j) - PREFIX(i)$$

¹<https://pdfs.semanticscholar.org/841a/31780b7e8f4de224fac06181321ca2ea807e.pdf>

For the static case, we can just precompute the answers! Directly store $B[i] := \sum_{j \leq i} A[j]$. On query, $B[i] = \text{PREFIX}(i)$ and thus we can answer in constant time. For example, if A is the following:

1	0	1	0	0	1	0	0	1	1
---	---	---	---	---	---	---	---	---	---

Then B would be:


1	1	2	2	2	3	3	3	4	5
---	---	---	---	---	---	---	---	---	---

Observation 1. *For the static case using $O(n)$ space, we can just precompute the answers for each key. We can then use a $O(\log \log n)$ predecessor search to get the partial sum up to that key.*

The following question remains: how do we maintain this structure when inserting new points?

2.2 Dynamic Case

The first idea is to keep a tree where the leaves point at the array A . Each node in the tree keeps track of the partial sums in its left and right subtrees. Thus, when we insert a new point x , we just update each node n along the path from root to leaf:



Theorem 2. *There exists a dynamic partial sum data structure with:*

$$t_q = O(\log n), t_u = O(\sqrt{\log n})$$

Main ideas: (1) Delay updates by buffering. (2) Exploit the self-reducibility of PS_n

Claim 3. *Suppose there exists a data structure \mathcal{D}_L for PS_{2^L} (a smaller array). Suppose it has update and query time t_u^L, t_q^L . Then we can design a data structure for PS_n using:*

$$t_u = O(t_u^L \cdot \frac{\log n}{L}), t_q = O(t_q^L \cdot \frac{\log n}{L})$$

Proof. We maintain a tree as before, but each node has fan out of size 2^L . Then the total height of the tree is $\frac{\log n}{L}$. Each node also maintains the smaller data structure \mathcal{D}_L .

- $\text{INSERT}(x)$: At each node, insert x into small partial sum data structures \mathcal{D}_L .
- $\text{QUERY}(x)$: Traverse tree, query each node for partial sums in children in time $\frac{\log n}{L}$.


The time to query and update each node is just t_q^L and t_u^L , thus the total query and update time is:

$$t_u = O(t_u^L \cdot \frac{\log n}{L}), t_q = O(t_q^L \cdot \frac{\log n}{L})$$

□

2.3 Buffer Trees

Note that there is room for improvement in our partial sum tree data structure. Our word size is $\log n$ and key size is L , but we don't exploit the fact that we can store lots of keys in a single word. When we perform updates, we only update a single bit as we traverse the tree. However, we can improve on this with Buffer Trees:



- Keep a buffer at each node of size $w = \log n$. \forall nodes $\in \mathcal{D}_L$, each node has buffer of most recent $\Theta(\frac{w}{L})$ updates. If buffer is not full, just insert into the buffer.
- *INSERT*(x):
 - (1) If current node buffer is not full, insert x into the buffer (do not reflect this new key in the partial sums of the left and right subtrees).
 - (2) If current node buffer is full, flush the buffer and distribute the updates to the children. Re-compute partial sums of both children ($O(1)$). Recurse if necessary.
- *QUERY*(x):

Traverse the tree and collect the partial sums + buffers at each node. Traversing the tree takes time L , thus total time is:

$$t_q = O(t_q^L \cdot \frac{\log n}{L}) = O(L \cdot \frac{\log n}{L}) = O(\log n)$$

Amortized insertion analysis:

$$\text{Cost}(t_u^L) = O(1) + \text{Amortized}(\text{Cost of flushing})$$

The cost of flushing a buffer is $O(1)$ and the buffer flushes only with $O(\frac{L}{w})$ frequency. A single inserted key cannot trigger more than L flushes total when going down the tree. Thus the amortized cost is: $O(\frac{L^2}{w})$:

$$\text{Cost}(t_u^L) = O(1) + O(\frac{L^2}{w})$$

To calculate the overall amortized cost in the original tree:

$$t_u = \frac{\log n}{L} (O(1) + O(\frac{L^2}{w})) = \frac{\log n}{L} + \frac{L \log n}{w}$$

Choosing $L = \sqrt{w}$ and using $w = O(\log n)$:

$$t_u = O(\sqrt{\log n})$$

OPEN: Is this optimal? Conjecture that $t_u = o(\sqrt{\log n}) \implies t_q = \omega(\log n)$

3 Lower Bounds: Chronogram Method

Theorem 4 (FS '89). *For all dynamic data structures for PS_n , $t_q \geq \Omega(\log_{t_u} n)$. This implies that:*

$$\max\{t_u, t_q\} \geq \Omega(\frac{\log n}{\log \log n})$$

Idea 1: Do a series of random insertions $\in_R \{0, 1\}$ into random locations of array A , and then perform a random query. After n random updates to A , a random PS_n query $q \in_R [n]$ must read a lot of memory cells.

Idea 2: Divide n random updates into geometrically decaying epochs $U_k \cdots U_1$:

$$U_k = \boxed{u_1} \boxed{u_2} \boxed{u_3} \cdots \boxed{} \boxed{} \cdots U_i = \boxed{} \boxed{} \boxed{} \cdots U_1 = \boxed{}$$

In each epoch $|U_i| = \beta^i$, where $\beta = (t_u \cdot w)^3$ and $k = \Theta(\log_\beta n)$. We then insert β^i random updates into evenly spaced locations:

$$\forall j = 1 \dots \beta^i, A[j \cdot \frac{n}{\beta^i}] := u_j$$

where $u_j \in_R \{0, 1\}$. For example:

- U_1 updates $A[0 \frac{n}{\beta}], A[1 \frac{n}{\beta}], A[2 \frac{n}{\beta}] \cdots$
- U_2 updates $A[0 \frac{n}{\beta^2}], A[1 \frac{n}{\beta^2}], A[2 \frac{n}{\beta^2}] \cdots$
- U_k updates $A[0 \frac{n}{\beta^k}], A[1 \frac{n}{\beta^k}], A[2 \frac{n}{\beta^k}] \cdots$

etc. Remember that the updates are processed from $U_k \rightarrow U_1$.

Claim 5. *Geometric decay reduces a dynamic problem on n updates to roughly $\log n$ independent state problems.*

Claim 6. *The only memory cells in the data structure that reveal substantial information about U_i are cells written during that epoch.*

Let $D(U_i)$ be the memory state of the data structure after epoch U_i . Let $A_i :=$ the set of memory cells last written during U_i . This is equivalent to a partition of the memory state into $\log n$ colors. For example if we denote A_i as red (r), A_{i-1} as blue (b), A_{i-2} as green (g):

$$D(U_{i-2}) = \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline r & b & r & b & r & r & r & b & g & r \\ \hline \end{array}$$

The idea is that a certain number of registers A_i for epoch U_i must be queried to reflect the events from epoch U_i . To show this, consider how many bits of information about U_i can be revealed by the past and future epochs:

- Past $A_{>i}$: These reveal no information about U_i because past updates are independent in that they happened beforehand.
- Future $A_{<i}$: These are not necessarily independent from U_i . Its possible a memory cell in the future copied some memory cell that was written during the epoch U_i . But considering that the number of updates decays geometrically, very few cells should be written in the future.

Calculating the number of cells that can be written after U_i :

$$\begin{aligned} \sum_{j=1}^{i-1} |U_j| \cdot t_u \cdot w &= \sum_{j=1}^{i-1} \beta^j (t_u w) \\ &\leq 5\beta^{i-1} \cdot t_u w && \text{Since } \beta_j \text{ is decaying} \\ &<< \beta_i = |U_i| && \text{Since } \beta_i = (t_u \cdot w)^3 \end{aligned}$$

Lemma 7. *For large epochs (any epoch with size $>$ a small constant), we have that:*

$$\mathbb{E}_{q,U}[|D(q) \cap A_i|] \geq \Omega(1)$$

when $D(q)$ reads t_q memory cells on query q

Note that this implies the total size of the data structure over random query is at least:

$$\begin{aligned} \mathbb{E}_{q,U}[|D(q)|] &\geq \sum_{i=1}^k \mathbb{E}[|D(q) \cap A_i|] && A_i \text{ are disjoint} \\ &= \sum_{i=1}^k \Omega(1) \\ &= \Omega(\log_\beta n) \end{aligned}$$

Proof. Assume for purpose of contradiction that the Lemma is false, and there is an epoch where $\mathbb{E}_{q,U}[|D(q) \cap A_i|] = o(1)$. Let epoch U_i be this epoch. Then 99% of partial sum queries $q \in [n]$ do not read cells A_i . Consider all other epochs fixed. Alice's input is all epochs $U_k \cdots U_1$, while Bob's input is all epochs except for U_i . We construct an impossible compression scheme such that we can encode β^i random updates in $< \beta^i$ bits.

	Alice	Bob
Input	$U_k, U_{k-1}, \dots, U_i, \dots U_1$	$U_k, U_{k-1}, \dots, (?), \dots U_1$

Idea 1: Alice sends Bob all updated contents $A_{<i} = o(\beta_i)$.

Idea 2: Alice sends parity $\in \{0, 1\}$ for 1% of queries that touch A_i .

Decoding: Bob simulates his data structure for epochs $U_{i+1} \cdots U_k$ to get cells $A_{>i}$. Bob then updates his data structure with the contents of $A_{<i}$ from Alice. The only cells Bob doesn't know are A_i . Bob uses parity of queries from Alice (for 1% of queries that touch A_i) and existing data structure to reconstruct the partial sums for any new query.

Complexity: The size of Alice's first message is the number of cells written by epochs $U_1 \cdots U_{i-1}$ which can be encoded in $\leq \frac{\beta_i}{4}$ bits. The size of Alice's parity messages are:

$$\log \left(\frac{\beta_i}{\beta_i/100} \right) \approx \beta_i \frac{\log 100}{100} < \frac{\beta_i}{4}$$

Thus, the total size of Alice's messages is $\frac{\beta_i}{4} + \frac{\beta_i}{4} < \beta_i$. This is a contradiction, as the encoding should be at least β_i .

□